

# Task Graph Executor - API Documentation

## Table of Contents

1. [Core Classes](#)
  2. [API Reference](#)
  3. [Configuration](#)
  4. [Examples](#)
  5. [Error Handling](#)
  6. [Best Practices](#)
- 

## Core Classes

### Task

Represents a single executable unit with dependencies.

```
@dataclass
class Task:
    name: str          # Unique identifier
    func: Callable      # Function to execute
    dependencies: List[str] = []  # Task names this depends on
    args: tuple = ()    # Positional arguments
    kwargs: dict = {}   # Keyword arguments
    retries: int = 0     # Retry attempts on failure
    timeout: Optional[float] = None # Execution timeout
```

### Example:

```
task = Task(
    name="process_data",
    func=process_function,
    dependencies=["fetch_data"],
    args=(input_data,),
    kwargs={"mode": "fast"},
    retries=3
)
```

## TaskGraphExecutor

Main executor class that orchestrates task execution.

```
executor = TaskGraphExecutor(max_workers=4)
```

### Parameters:

- `max_workers` (int): Maximum number of concurrent tasks (default: 4)
- 

# API Reference

## Adding Tasks

### `add_task(task: Task)`

Add a task to the execution graph.

```
executor.add_task(Task("compile", compile_code))
executor.add_task(Task("test", run_tests, dependencies=["compile"]))
```

#### Raises:

- `ValueError`: If task with same name already exists

## Validation

### `validate_graph()`

Validate the task graph for cycles and missing dependencies.

```
executor.validate_graph()
```

#### Raises:

- `CyclicDependencyError`: If circular dependency detected
- `ValueError`: If dependency references non-existent task

**Implementation:** Uses Depth-First Search (DFS) with recursion stack to detect cycles in  $O(V + E)$  time.

## Execution

### `execute() -> Dict[str, Any]`

Execute all tasks in the graph.

```
results = executor.execute()
# Returns: {"task_name": result, ...}
```

#### Behavior:

1. Validates graph for cycles
2. Executes tasks in topological order
3. Maximizes parallelism by running independent tasks simultaneously
4. Handles failures and retries
5. Returns results dictionary

**Returns:** Dictionary mapping task names to their return values

## Topological Sorting

`topological_sort() -> List[str]`

Returns tasks in dependency order.

```
order = executor.topological_sort()  
# Returns: ["task_a", "task_b", "task_c"]
```

**Implementation:** Kahn's algorithm with in-degree tracking

## Statistics

`get_execution_stats() -> Dict[str, Any]`

Get detailed execution statistics.

```
stats = executor.get_execution_stats()
```

**Returns:**

```
{  
    "total_tasks": int,  
    "completed": int,  
    "failed": int,  
    "running": int,  
    "task_details": {  
        "task_name": {  
            "status": str,  
            "duration": float,  
            "retries": int,  
            "error": Optional[str]  
        }  
    }  
}
```

## Visualization

`visualize_graph() -> str`

Generate ASCII representation of task graph.

```
print(executor.visualize_graph())
```

**Output:**

Task Dependency Graph:

```
=====  
compile_main (no dependencies)  
compile_utils (no dependencies)  
link <- ['compile_main', 'compile_utils']
```

```
test <- ['link']
```

---

## Configuration

### Thread Pool Sizing

```
# CPU-bound tasks
executor = TaskGraphExecutor(max_workers=cpu_count())

# I/O-bound tasks
executor = TaskGraphExecutor(max_workers=cpu_count() * 2)

# Mixed workload
executor = TaskGraphExecutor(max_workers=8)
```

### Retry Configuration

```
task = Task(
    "flaky_api_call",
    api_function,
    retries=3 # Will attempt up to 4 times total
)
```

### Retry Behavior:

- Immediate retry (1 second delay)
  - Exponential backoff not implemented (can be added)
  - Final failure marks task as FAILED
- 

## Examples

### Example 1: Build System

```
executor = TaskGraphExecutor(max_workers=4)

# Compile modules in parallel
executor.add_task(Task("compile_main", compile, args=("main.c",)))
executor.add_task(Task("compile_utils", compile, args=("utils.c",)))
executor.add_task(Task("compile_config", compile, args=("config.c",)))

# Link after all compilation completes
executor.add_task(Task(
    "link",
    link_objects,
    dependencies=["compile_main", "compile_utils", "compile_config"]
))

# Test the binary
```

```
executor.add_task(Task("test", run_tests, dependencies=["link"]))
```

```
results = executor.execute()
```

## Example 2: Data Pipeline

```
executor = TaskGraphExecutor(max_workers=3)
```

```
# Extract from multiple sources
```

```
executor.add_task(Task("extract_users", extract_data, args=("users_db",)))
executor.add_task(Task("extract_orders", extract_data, args=("orders_db",)))
```

```
# Transform extracted data
```

```
executor.add_task(Task(
    "transform_users",
    transform,
    args=("normalize",),
    dependencies=["extract_users"]
))
```

```
executor.add_task(Task(
    "transform_orders",
    transform,
    args=("aggregate",),
    dependencies=["extract_orders"]
))
```

```
# Load to warehouse
```

```
executor.add_task(Task(
    "load_warehouse",
    load_data,
    dependencies=["transform_users", "transform_orders"]
))
```

```
results = executor.execute()
```

## Example 3: With Visualization

```
from task_graph_visualization import VisualizationExecutor
```

```
executor = VisualizationExecutor(max_workers=4, monitoring_port=8090)
```

```
# Add tasks...
```

```
executor.add_task(Task("step1", process))
executor.add_task(Task("step2", validate, dependencies=["step1"]))
```

```
# Start monitoring server
```

```
executor.start_monitoring()
print("Dashboard: http://localhost:8090")
```

```
# Execute with live monitoring
```

```
results = executor.execute()
```

```
# Keep server running
input("Press Enter to stop...")
executor.stop_monitoring()
```

---

## Error Handling

### Task Failures

When a task fails:

1. Task status set to FAILED
2. Error captured in `task.error`
3. Dependent tasks marked as SKIPPED
4. Execution continues for independent tasks

```
try:
    results = executor.execute()
except Exception:
    # Graph validation failed
    pass

# Check individual task failures
stats = executor.get_execution_stats()
for name, details in stats['task_details'].items():
    if details['status'] == 'FAILED':
        print(f"{name} failed: {details['error']}")
```

### Common Exceptions

| Exception                          | Cause                                | Solution                           |
|------------------------------------|--------------------------------------|------------------------------------|
| <code>CyclicDependencyError</code> | Circular dependencies detected       | Review task dependencies           |
| <code>ValueError</code>            | Missing dependency or duplicate name | Check task names and dependencies  |
| <code>Exception</code> (from task) | Task execution failed                | Check task function implementation |

## Debugging Failed Executions

```
# Get detailed stats
stats = executor.get_execution_stats()

# Check failed tasks
failed = [name for name, details in stats['task_details'].items()
          if details['status'] == 'FAILED']

# Inspect specific task
task = executor.tasks['failed_task_name']
print(f"Error: {task.error}")
print(f"Retries: {task.retry_count}")
print(f"Duration: {task.duration()}")
```

---

# Best Practices

## 1. Task Granularity

**Good:** Fine-grained tasks enable better parallelism

```
executor.add_task(Task("download", download_file))
executor.add_task(Task("parse", parse_file, dependencies=["download"]))
executor.add_task(Task("validate", validate, dependencies=["parse"]))
```

**Bad:** Coarse-grained tasks limit parallelism

```
executor.add_task(Task("do_everything", download_parse_validate))
```

## 2. Dependency Management

**Good:** Explicit, minimal dependencies

```
# Task C only needs A, not B
executor.add_task(Task("C", func_c, dependencies=["A"]))
```

**Bad:** Unnecessary dependencies

```
# Task C depends on B unnecessarily
executor.add_task(Task("C", func_c, dependencies=["A", "B"]))
```

## 3. Worker Configuration

```
# For CPU-bound tasks (computation)
max_workers = os.cpu_count()
```

```
# For I/O-bound tasks (network, disk)
max_workers = os.cpu_count() * 2
```

```
# For mixed workloads
max_workers = os.cpu_count() + 2
```

## 4. Error Handling

```
def robust_task():
    try:
        result = risky_operation()
        return result
    except SpecificError as e:
        # Handle expected errors
        return default_value
    # Let unexpected errors propagate for retry
```

## 5. Idempotency

Design tasks to be idempotent (can be safely retried):

```
def idempotent_task(output_path):
    # Check if already done
    if os.path.exists(output_path):
        return load_existing(output_path)

    # Do work
    result = process_data()
    save_result(result, output_path)
    return result
```

## 6. Logging

```
import logging
```

```
def logged_task(data):
    logging.info(f"Processing {len(data)} items")
    result = process(data)
    logging.info(f"Completed with {len(result)} results")
    return result
```

---

# Performance Considerations

## Time Complexity

| Operation      | Complexity | Notes         |
|----------------|------------|---------------|
| Add Task       | O(1)       | Amortized     |
| Validate Graph | O(V + E)   | DFS traversal |

|                  |            |                                 |
|------------------|------------|---------------------------------|
| Topological Sort | $O(V + E)$ | Kahn's algorithm                |
| Get Ready Tasks  | $O(V)$     | Check all tasks                 |
| Execute          | $O(V + E)$ | Plus actual task execution time |

Where  $V$  = number of tasks,  $E$  = number of dependencies

## Space Complexity

- Task storage:  $O(V)$
- Dependency graphs:  $O(V + E)$
- Thread pool:  $O(\text{max\_workers})$
- Total:  $O(V + E + \text{max\_workers})$

## Optimization Tips

1. **Minimize dependency chains:** Deeper chains reduce parallelism
  2. **Balance worker count:** Too few = underutilized, too many = overhead
  3. **Use retries sparingly:** Only for transient failures
  4. **Profile task duration:** Identify bottlenecks in critical path
- 

## Thread Safety

The executor is thread-safe for:

- Adding tasks (before execution)
- Concurrent task execution
- Statistics retrieval

Not thread-safe for:

- Adding tasks during execution
- Modifying task properties during execution

# Safe

```
executor.add_task(task1)
executor.add_task(task2)
results = executor.execute()
```

# Unsafe

```
def background_add():
    executor.add_task(new_task) # Don't do this during execution

executor.execute()
```

---

## Limitations

1. **No distributed execution:** Single-machine only
2. **No task priority:** All tasks at same level have equal priority
3. **No dynamic dependencies:** Dependencies must be known upfront
4. **No resource constraints:** Doesn't track memory, CPU, etc.
5. **Simple retry:** No exponential backoff or jitter

## Future Enhancements

Potential improvements for production use:

- Distributed execution across multiple machines
- Priority-based scheduling
- Resource constraints (memory, CPU limits)
- Advanced retry strategies (exponential backoff)
- Persistent state for resuming failed executions
- Task result caching
- Dynamic task generation
- Integration with job queues (Celery, RQ)

---

## Appendix: Algorithm Details

### Cycle Detection (DFS)

```
def has_cycle(node, visited, recursion_stack):  
    visited.add(node)  
    recursion_stack.add(node)  
  
    for neighbor in dependencies[node]:  
        if neighbor not in visited:  
            if has_cycle(neighbor, visited, recursion_stack):  
                return True  
        elif neighbor in recursion_stack:  
            return True # Back edge found = cycle  
  
    recursion_stack.remove(node)  
    return False
```

### Topological Sort (Kahn's Algorithm)

```
def topological_sort():  
    in_degree = {task: len(dependencies[task]) for task in tasks}  
    queue = [task for task, degree in in_degree.items() if degree == 0]  
    sorted_tasks = []  
  
    while queue:  
        task = queue.pop(0)
```

```
sorted_tasks.append(task)

for dependent in dependents[task]:
    in_degree[dependent] -= 1
    if in_degree[dependent] == 0:
        queue.append(dependent)

return sorted_tasks
```

---

For additional examples and use cases, see the test suite in [task\\_graph\\_tests.py](#).